

Implementation and Analysis of a Knowledge-Based Chatbot in C++ Practical

Report September 11, 2025 Prepared by: Grok (xAI Assistant)

1 Title

Implementation and Analysis of a Knowledge-Based Chatbot in C++

2 Objective

The objective of this practical is to develop a C++ program that implements a simple chatbot, named SmartBot, capable of:

- Responding to user queries based on a stored knowledge base.
- Learning new question-answer pairs interactively from the user.
- Persisting the knowledge base to a file for future use.
- Demonstrating string manipulation, file I/O, and data structures like unordered maps.

This practical aims to showcase basic natural language interaction, persistent storage, and the use of C++ standard libraries for building an extensible conversational system.

3 Introduction

Chatbots are software agents designed to simulate human conversation, often used in customer service, education, or entertainment. This implementation creates a console-based chatbot that uses a knowledge base stored in a text file to answer user queries. The bot supports case-insensitive matching, learns new responses, and saves them for persistence. The program leverages C++ standard libraries such as `<iostream>`, `<string>`, `<unordered_map>`, `<fstream>`, and `<algorithm>` for input/output, string processing, hash table storage, file handling, and text transformation, respectively.

The N-Queens problem, discussed in a previous practical, used backtracking to solve a combinatorial puzzle. In contrast, this chatbot focuses on interactive data management and string processing, highlighting different aspects of algorithmic design and practical application in C++.

4 Methodology (What Was Done)

The program was developed with the following steps:

1. Initialization: Load an existing knowledge base from a text file (knowledge.txt) into an unordered map, where keys are questions (lowercase) and values are answers.
2. User Interaction: Prompt the user for input in a loop until they type "exit".
3. Query Processing: Convert user input to lowercase and check the knowledge base for a matching question.
4. Response Handling: If a match is found, display the stored answer. Otherwise, offer to learn a new answer, which is then saved to the knowledge base and file.

5. Persistence: Save new question-answer pairs to the file to ensure persistence across sessions.

The program was tested with sample inputs, including known and unknown questions, to verify functionality and file operations.

5 Implementation Details (How It Was Done)

The code is modular, with functions for string processing, file I/O, and the main interaction loop.

5.1 toLowerCase(string str)

- Purpose: Converts a string to lowercase for case-insensitive matching.
- How It Works: Uses `std::transform` with `::tolower` to convert each character.
- Efficiency: $O(n)$, where n is the string length.

5.2 loadKnowledgeBase(const string& filename)

- Purpose: Loads question-answer pairs from a file into an unordered map.
- How It Works: Reads the file line by line, treating alternating lines as questions and answers, converting questions to lowercase.
- Key Aspects: Returns an empty map if the file doesn't exist. Assumes correct file format (question and answer pairs).

5.3 saveKnowledgeBase(const unordered_map<string, string>& kb, const string& filename)

- Purpose: Saves the knowledge base to a file.
- How It Works: Writes each question-answer pair to the file, separated by newlines.
- Key Aspects: Overwrites the file, ensuring the latest knowledge base is stored.

5.4 main()

- Purpose: Manages user interaction and program flow.
- How It Works:
 - Loads the knowledge base.
 - Enters a loop prompting for user input.
 - Exits if input is "exit".
 - Checks the knowledge base for the lowercase input.
 - If no match, asks to teach a new answer, updating the knowledge base and file if provided.

5.5 Compilation and Execution

- Compiled with a C++ compiler (e.g., g++ with C++11 or later).
- Example Run: User inputs "hello", gets a response if known, or teaches a new response like "Hi there!" which is saved.

6 Results

The program was tested with a sample knowledge.txt file containing:

what is the capital of france

Paris

who is elon musk

CEO of Tesla and SpaceX

Sample interaction:

SmartBot: Hello! I am your assistant. Type 'exit' to quit.

You: what is the capital of france

SmartBot: Paris

You: what is python

SmartBot: I dont know the answer. Can you teach me? (yes/no)

You: yes

Please provide the correct answer:

A programming language

SmartBot: Thank you! I will remember this.

You: exit

SmartBot: Goodbye!

After teaching, knowledge.txt updated to include:

what is python

A programming language

The program correctly loaded, responded, learned, and saved new knowledge, with no crashes or errors.

7 Analysis in Detail

7.1 Algorithm Correctness

The chatbot correctly handles:

- Case-insensitive queries using toLowerCase.
- File I/O for persistent storage.
- Interactive learning, updating the knowledge base in memory and on disk.

Correctness was verified by matching known responses and checking file updates.

7.2 Time Complexity

- toLowerCase: $O(n)$, where n is string length.
- loadKnowledgeBase: $O(m)$, where m is the number of lines in the file.
- saveKnowledgeBase: $O(k)$, where k is the number of entries in the knowledge base.
- Lookup in unordered_map: $O(1)$ average case.
- Main loop: Depends on user interactions, but each query is $O(n)$ for string conversion plus $O(1)$ for lookup.

7.3 Space Complexity

- $O(k)$ for the unordered_map, where k is the number of question-answer pairs.
- $O(n)$ for temporary strings during input processing.

7.4 Performance Observations

- Responsiveness: Instantaneous for small knowledge bases due to $O(1)$ hash table lookups.
- File I/O: Bottleneck for large files, but negligible for typical use (e.g., <1000 entries).
- Limitations: No error handling for file issues (e.g., missing or corrupted knowledge.txt). Assumes well-formed input. No partial matching or fuzzy search for queries.
- Scalability: Suitable for small to medium knowledge bases. Large bases may slow down file I/O but not lookups.

7.5 Strengths

- Simple and extensible design.
- Persistent storage ensures knowledge retention.
- Case-insensitive matching improves usability.

7.6 Weaknesses

- No error handling for invalid file formats or I/O failures.
- Exact-match queries limit flexibility (e.g., no synonyms or partial matches).
- Console-based; no GUI for better user experience.

8 Conclusion

This practical successfully implemented a knowledge-based chatbot in C++, demonstrating file I/O, hash table usage, and interactive learning. The program is efficient for small-scale use but could be enhanced with error handling, partial matching, or a graphical interface. Future improvements might include natural language processing for smarter query matching or database integration for larger knowledge bases.

9 References

- C++ Standard Library documentation for `unordered_map` and `fstream`.
- "The C++ Programming Language" by Bjarne Stroustrup.